

CS6140 Machine Learning Fall 2014 Homework 3, Wei Luo

PROBLEM 1

Using Gaussian Discriminant Analysis, I got:

average train error rate: 0.091212 average test error rate: 0.096498

Since the average accuracy is over 0.9, it seems to me that the gaussian assumption holds for this data set.

PROBLEM 2

The Error Table for Naïve Bayes Classifier, Model with Bernoulli Random Variables:

fold#	false positive rate	false negative rate	overall error rate
1	0.142322	0.067010	0.110629
2	0.134021	0.112426	0.126087
3	0.123288	0.089286	0.110870
4	0.104651	0.084158	0.095652
5	0.119718	0.062500	0.097826
6	0.127208	0.056497	0.100000
7	0.118467	0.127168	0.121739
8	0.114695	0.049724	0.089130
9	0.128571	0.083333	0.110870
10	0.123596	0.062176	0.097826
average	0.12365375	0.07942785	0.10606291

The Error Table for Naïve Bayes Classifier, Model with Gaussian Random Variables:

fold#	false positive rate	false negative rate	overall error rate
1	0.186207	0.140351	0.169197
2	0.121622	0.091463	0.110870
3	0.102273	0.056122	0.082609
4	0.138996	0.084577	0.115217
5	0.175000	0.077778	0.136957
6	0.178832	0.112903	0.152174
7	0.136842	0.085714	0.117391
8	0.138686	0.096774	0.121739
9	0.199301	0.137931	0.176087
10	0.157143	0.094444	0.132609
average	0.15349013	0.09780588	0.13148496

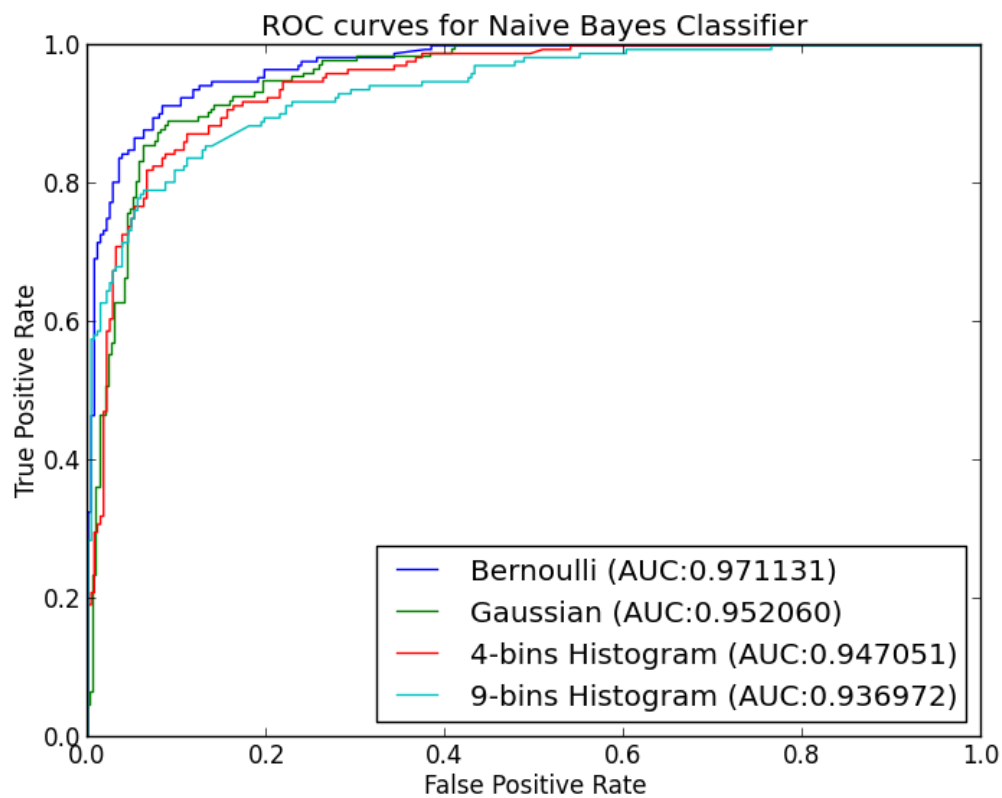
The Error Table for Naïve Bayes Classifier (4-bins Histogram):

fold#	false positive rate	false negative rate	overall error rate
1	0.086957	0.221622	0.140998
2	0.097222	0.151163	0.117391
3	0.076642	0.139785	0.102174
4	0.103321	0.190476	0.139130
5	0.064982	0.218579	0.126087
6	0.084559	0.196809	0.130435
7	0.078292	0.195531	0.123913
8	0.081851	0.201117	0.128261
9	0.094737	0.194286	0.132609
10	0.063604	0.225989	0.126087
average	0.08321663	0.19353558	0.12670848

The Error Table for Naïve Bayes Classifier (9-bins Histogram):

fold#	false positive rate	false negative rate	overall error rate
1	0.032374	0.251366	0.119306
2	0.025180	0.340659	0.150000
3	0.025271	0.371585	0.163043
4	0.007692	0.485000	0.215217
5	0.021818	0.356757	0.156522
6	0.027119	0.400000	0.160870
7	0.027027	0.358209	0.171739
8	0.049123	0.291429	0.141304
9	0.023649	0.329268	0.132609
10	0.017544	0.365714	0.150000
average	0.02567962	0.3549987	0.15606102

The ROC curves and AUC for each classifier is:



```

function [label, model, llh] = emgm(X, init)
% Perform EM algorithm for fitting the Gaussian mixture model.
% X: d x n data matrix
% init: k (1 x 1) or label (1 x n, 1 ≤ label(i) ≤ k) or center (d x k)
% Written by Michael Chen (sth4nth@gmail.com).
%% initialization
fprintf('EM for Gaussian mixture: running ... \n');
R = initialization(X, init);
[~, label(1,:)] = max(R, [], 2);
R = R(:, unique(label));

tol = 1e-10; converge threshold 10-10
maxiter = 500; maximum 500 steps
llh = -inf(1, maxiter); initialize loglikelihood with -inf
converged = false; init converged
t = 1; init t
while ~converged && t < maxiter loop while not converged and not reach maxiter
    t = t + 1;
    model = maximization(X, R); m-step
    [R, llh(t)] = expectation(X, model); e-step

    [~, label(:)] = max(R, [], 2);
    u = unique(label); % non-empty components
    if size(R, 2) ~= size(u, 2)
        R = R(:, u); % remove empty components
    else
        converged = llh(t) - llh(t-1) < tol * abs(llh(t)); converged if difference between current and previous likelihood less than threshold
    end
end
llh = llh(2:t); llh value starts from llh(2), first value not set in loop
if converged converged
    fprintf('Converged in %d steps.\n', t-1);
else not converged in maxiter steps
    fprintf('Not converged in %d steps.\n', maxiter);
end

function R = initialization(X, init)
[d, n] = size(X); get X (d x n data) dimensions
if isstruct(init) % initialize with a model
    R = expectation(X, init); return expectation for X from given model
elseif length(init) == 1 % random initialization
    k = init;
    idx = randsample(n, k); randomly pick sample k indexes from 1 to n
    m = X(:, idx); get sample data
    [~, label] = max(bsxfun(@minus, m'*X, dot(m, m, 1)'/2), [], 1);
    [u, ~, label] = unique(label);
    while k ~= length(u)
        idx = randsample(n, k); pick random samples
        m = X(:, idx);
        [~, label] = max(bsxfun(@minus, m'*X, dot(m, m, 1)'/2), [], 1);
        [u, ~, label] = unique(label);
    end
    R = full(sparse(1:n, label, 1, n, k, n)); make a full sparse matrix that is a random membership matrix
elseif size(init, 1) == 1 && size(init, 2) == n % initialize with labels
    label = init;
    k = max(label);
    R = full(sparse(1:n, label, 1, n, k, n));
elseif size(init, 1) == d % initialize with only centers
    k = size(init, 2);
    m = init;
    [~, label] = max(bsxfun(@minus, m'*X, dot(m, m, 1)'/2), [], 1);
    R = full(sparse(1:n, label, 1, n, k, n));
else
    otherwise, init is not valid.

```

```

error('ERROR: init is not valid.');
```

end

```

function [R, llh] = expectation(X, model)
mu = model.mu;
Sigma = model.Sigma;
w = model.weight;
n = size(X,2);
k = size(mu,2);
logRho = zeros(n,k);
for i = 1:k
    logRho(:,i) = loggausspdf(X,mu(:,i),Sigma(:,i));
end
logRho = bsxfun(@plus,logRho,log(w));
T = logsumexp(logRho,2);
llh = sum(T)/n; % loglikelihood
logR = bsxfun(@minus,logRho,T);
R = exp(logR);
function model = maximization(X, R)
[d,n] = size(X);
k = size(R,2);
nk = sum(R,1);
w = nk/n;
mu = bsxfun(@times, X*R, 1./nk);
Sigma = zeros(d,d,k);
sqrtR = sqrt(R);
for i = 1:k
    Xo = bsxfun(@minus,X,mu(:,i));
    Xo = bsxfun(@times,Xo,sqrtR(:,i));
    Sigma(:,i,i) = Xo*Xo'/nk(i);
    Sigma(:,i,i) = Sigma(:,i,i)+eye(d)*(1e-6); % add a prior for numerical stability
end
model.mu = mu;
model.Sigma = Sigma;
model.weight = w;
function y = loggausspdf(X, mu, Sigma)
d = size(X,1);
X = bsxfun(@minus,X,mu);
[U,p]= chol(Sigma);
if p ~= 0
    error('ERROR: Sigma is not PD.');
```

end

```

Q = U'\X;
q = dot(Q,Q,1); % quadratic term (M distance)
c = d*log(2*pi)+2*sum(log(diag(U))); % normalization constant
y = -(c+q)/2;
```

PROBLEM 4

A) The results for 2 gaussian:

Gaussian₁: points: 1981; mean: [2.94846981, 3.0421113]
covariance: [[0.8992528, 0.01219505], [0.01219505, 2.96950283]]

Gaussian₂: points: 4019; mean: [7.0078808, 3.98162176]
covariance: [[0.95936573, 0.4865028], [0.4865028, 0.99568567]]

B) The results for 3 gaussian:

Gaussian₁: points: 3098; mean: [6.98857865, 4.03307065]
covariance: [[1.00480526, 0.49076192], [0.49076192, 1.0080817]]

Gaussian₂: points: 1945; mean: [2.91656552, 2.8932281]
covariance: [[0.86404608, -0.20507606], [-0.20507606, 2.9738529]]

Gaussian₃: points: 4957; mean: [4.99043274, 7.03134783]
covariance: [[0.91578388, 0.18253052], [0.18253052, 0.90038684]]

PROBLEM 5

a)

$$\frac{P(B|A,C)P(A|C)}{P(B|C)} = \frac{P(A,B,C)/P(A,C) * P(A,C)/P(C)}{P(B,C)/P(C)} = \frac{P(A,B,C)}{P(B,C)} = P(A|B,C)$$

b) The prior probability of the coin being fair is $P(F) = \frac{F}{F+1}$. So the prior probability of the coin being double-headed is $P(D) = 1 - P(F) = \frac{1}{F+1}$. If we see n heads in a row, the probability that the coin is fair is: $P(F|nH) = \frac{P(nH|F)P(F)}{P(nH)}$; the probability that the coin is double-headed is: $P(D|nH) = \frac{P(nH|D)P(D)}{P(nH)}$. If the coin is fair, the probability of seeing a head is $P(H|F) = \frac{1}{2}$. the probability of seeing n heads in a row is $P(nH|F) = P(H|F)^n = (\frac{1}{2})^n$. If the coin is double-headed, the probability of seeing a head and n head in a row is $P(H|D) = P(nH|D) = 1$. We want a better than even chance that the coin is double-headed, which means $P(D|nH) > P(F|nH)$. So $\frac{P(nH|D)P(D)}{P(nH)} > \frac{P(nH|F)P(F)}{P(nH)}$. So $1 * \frac{1}{F+1} > (\frac{1}{2})^n * \frac{F}{F+1}$. We get $n > \log_2 F$. So we need to see more than $\log_2 F$ heads in a row to become convinced that there is a better than even chance that the coin is double-headed.